

Artificial Intelligence A-Z™: Additional Reading

Part 0: Fundamentals of Reinforcements Learning

Additional Reading:

- Arthur Juliani, 2016, [Simple Reinforcement Learning with Tensorflow \(10 Parts\)](#)
 - Richard Sutton et al., 1998, [Reinforcement Learning I: Introduction](#)
 - Richard Bellman, 1954, [The Theory of Dynamic Programming](#)
 - D. J. White, 1993, [A Survey of Applications of Markov Decision Processes](#)
 - Martijn van Otterlo, 2009, [Markov Decision Processes: Concepts and Algorithms](#)
 - Richard Sutton, 1988, [Learning to Predict by the Methods of Temporal Differences](#)
-

Part 1: Deep Q-Learning

Additional Reading:

- Arthur Juliani, 2016, [Simple Reinforcement Learning with Tensorflow \(Part 4\)](#)
 - Tom Schaul et al., Google DeepMind, 2016, [Prioritized Experience Replay](#)
 - Michel Tokic, 2010, [Adaptive \$\epsilon\$ -greedy Exploration in Reinforcement Learning Based on Value Differences](#)
-

Part 2: Deep Convolutional Q-Learning

Additional Reading:

- Richard S. Sutton and Andrew G. Barto, 1998, [Reinforcement Learning: An Introduction](#)
 - Volodymyr Mnih et al., 2016, [Asynchronous Methods for Deep Reinforcement Learning](#)
-

Part 3: Asynchronous Actor-Critic Agents (A3C)

Additional Reading:

- Volodymyr Mnih et al, 2016 [Asynchronous Methods for Deep Reinforcement Learning](#)
- Jaromír Janisch, 2017 [Let's Make An A3c: Implementation](#)
- John Schulman et al., 2016 [High-dimensional Continuous Control Using Generalized Advantage Estimation](#)
- Arthur Juliani, 2016 [Simple Reinforcement Learning with Tensorflow \(Part 8\)](#)